OS in 1000 Lines Notes

Katie Baek

Boston College

Notes from coding OS in 1000 lines

November 20, 2025

OS in 1000 Lines Notes

Katie Baek

November 20, 2025

Abstract

In the spirit of grokking the OS, we are attempting to write our own mini-OS using the guide "OS in 1000 Lines"[3]. We are expecting to finish within the next few weeks, then move onto the hypervisor. Unless otherwise stated, assume that the information is refering the guide.

Contents

1 RISC-V 101

The chosen CPU for this book is RISC-V, specifically, the 32-bit RISC-V. This is because there is well-written documentation and it has been rising in popularity in the recent years. The underlying machine is the QEMU virtual machine. Although this does not exist in the real world, it's simple and similar to a lot of real devices.

1.1 RISC-V Assembly

In order to write an OS, you need to write some assembly. Here is an example of assembly code:

```
1    addi a1 a0 123
```

The first instruction *addi* is the instruction name (opcode). The entire line tells the computer to add the value **123** to the value **a1** and store it in the register **a0**

1.2 Registers in RISC-V

Below is a table of common registers in RISC-V and their given alias that we will use to refer to them in code.

Register	ABI Name (alias)	Description
pc	pc	Program counter (where the next instruction is)
x0	zero	Hardwired zero (always reads as zero)
x1	ra	Return address
x2	sp	Stack pointer
x5–x7	t0–t2	Temporary registers
x8	fp	Stack frame pointer
x10–x11	a0–a1	Function arguments/return values
x12–x17	a2–a7	Function arguments
x18–x27	s0–s11	Temporary registers saved across calls
x28–x31	t3–t6	Temporary registers

1.3 Memory access

Most data is stored in memory, then moved to registers when its needed and so there are instructions which do this:

```
1    lw a0, (a1)
2    sw a0, (a1)
```

In line 1, you read a word from the address that is stored in `a1` and store it in `a0`. The equivalent C code is `a0 = *a1`. You are doing the same thing in line 2, but storing the word instead of loading it. The `(...)` is like a pointer dereference.

1.4 Stack

The stack is a LIFO (last in first out) memory space used for function calls and local variables. It grows downwards, and the stack pointer, `sp` points to the "top" of the stack (but functionally the bottom since again, it grows downwards). To push onto the stack, decrement the stack pointer and store the value:

```
1    addi sp, sp -4 // Move the stack pointer down by 4 bytes
2    sw a0, (sp) // Store a0 to the stack
```

To pop from the stack, load the value and increment the stack pointer:

```
1    lw a0, (sp) // Load a0 from the stack
2    addi sp, sp 4 // Move the stack pointer up by 4 bytes
```

1.5 Privileged instructions

With CPU instructions, there are privileged instructions that applications in user mode cannot execute. The **CSR (Control and Status Register)** is a register that sores CPU settings. We use the following:

Opcode and operands	Overview
`csrr rd, csr`	Read from CSR
`csrw csr, rs`	Write to CSR
`csrrw rd, csr, rs`	Read from and write to CSR at once
`sret`	Return from trap handler
`sfence.vma`	Clear Translation Lookaside Buffer (TLB)

1.6 Inline assembly

Inline assembly code for C looks like this:

```
1   uint32_t value;
2   __asm__ __volatile__("csrr %0, sepc" : "=r(value)");
```
Listing 1: sample inline assembly code

Where its written in the form of:

```
1   __asm__ __volatile__("assembly" : output operands : input operands : clobbered
        registers);
```
Listing 2: format of inline assembly code

Part	Description
`__asm__`	Indicates its inline assembly
`__volatile__`	Tell the compiler not to optimize the "assembly" code
`"assembly"`	Assembly code written as a string literal
output operands	C variables to store the results of the assembly
input operands	C expressions to be used in the assembly
clobbered registers	Registers whose contents are destroyed in the assembly. If forgotten, the C compiler won't preserve the contents of these registers and would cause a bug.

In the code 1, this reads the value of `sepc` CSR using the `scrr` instruction and assigns it to the `value` variable.

```
1   __asm__ __volatile__("csrw sscratch, %0" : : "r"(123))
```
Listing 3: different sample inline assembly code

In the code 3, this writes `123` to the `sscratch` CSR, using the `csrw` instruction. The C compiler automatically sets `123` to the register `a0` so that the assembly can be executed as written which is nice because then the dev doesn't have to worry about it.

2 Overview

2.1 Features

We will implement the following features:

1. **Multitasking**: Switch between processes to allow multiple applications to share the CPU

2. **Execption Hanlder**: Handle events requiring OS intervention, such as illegal instructions

3. **Paging**: Provide an isolated memory address space for each application

4. **System calls**: Allow applications to call kernel features

5. **Device drivers**: Abstract hardware functionalities, such as disk read/write

6. **File system**: Manage files on disk

7. **Command-line shell**: User interface for humans

3 Boot

What happens when the computer is turned on? The CPU initalizes itself and starts executing the OS. The OS then initalizes the hardware and starts the applications. This process is called "booting".

What happens before the OS starts? In PCs, BIOS initalizes the hardware, displays the splash screen, and loads the OS from the disk. In QEMU virt machine, **OpenSBI** is the equivalent of BIOS. This is the CPU "initalizing itself".

3.1 Supervisor Binary Interface (SBI)

The SBI is an API for OS kernels, but defines what the firmware (OpenSBI) provides to an OS. In QEMU, OpenSBI starts by default, performs hardware-specific initalization, and boots the kernel. In other words, the SBI performs tasks on the hardware/controls the hardware on behalf of the OS. It can be analogous to a program making sys calls to the kernel.

3.2 Booting OpenSBI

The following shell code starts OpenSPI:

```
1   #!bin/bash
2   set -xue
3
4   # QEMU file path
5   QEMU=qemu-system-riscv32
6
7   # Start QEMU
8   $QEMU -machine virt -bios default -nographic -serial mon:stdio --no-reboot
```
Listing 4: script to start OpenSBI

With this shell script, you can boot OpenSBI. Specifically, the `-bios default` option tells QEMU to use the default firmware, which is OpenSBI. In addition, at this point, QEMU's standard input and standard output is connected to the virtual machine's serial port, and the characters that you type in are being sent to OpenSBI but there is no one to read the characters. This is because there is literally no OS. In order to do anything, we have to interact with QEMU (which is the software emulator and a type-2 hypervisor).

Press Ctrl+A then C to switch to the QEMU debug console (QEMU monitor)

Here are the other options which are used in the run script:

1. `-machine vert`: Start a `vert` machine

2. `-bios default`: Use the default firmware (OpenSBI)

3. `-nographic`: Start QEMU without a GUI window

4. `-serial mon:stdio`: Connect QEMU's standard IO to the virtual machine's serial port. Specifying `mon:` allows switching to the QEMU monitor by pressing Ctrl + A then C

5. `--no-reboot`: If the virtual machine crashes, stop the emulator without rebooting (useful for debugging)

4

3.3   Linker script

```
1 ENTRY(boot)
2
3 SECTIONS {
4   . = 0x80200000;
5
6   .text :{
7     KEEP(*(.text.boot));
8     *(.text .text.*);
9   }
10
11   .rodata : ALIGN(4) {
12     *(.rodata .rodata.*);
13   }
14
15   .data : ALIGN(4) {
16     *(.data .data.*);
17   }
18
19   .bss : ALIGN(4) {
20     __bss = .;
21     *(.bss .bss.* .sbss. .sbss.*);
22     __bss_end = .;
23   }
24
25   . = ALIGN(4);
26   . += 128 * 1024;
27   __stack_top = .;
28 }
```

Listing 5: the linker script

The linker script is a file which defines the memory layout of executable files. Based on the layout, the linker assigns memory address to functions and variables. AKA this tells the linker exactly where in memory all parts of the program should live. You need this because during the **boot** process, there is no OS, no loader, and no runtime environment to relocate the code for you. Specifically, the linker script controls:

1. Where the code is placed in memory

2. Where global variables, stacks, and boot code lives

3. The address of symbols used by startup code

4. How the final binary is organized

3.3.1   Why do you need one during boot?

When a computer boots, the CPU starts executing from a well-defined hardware reset address. In this case, it is 0x80200000. The boot code *must* be there or else the CPU would start executing whatever is at that address. In addition, it also defines the variables __bss, __bss_end, and __stack_top with specific addresses.

The boot code needs to ensure that the hardware is in a state to find and load the kernel, which is then going to take things over from there [1]

3.4   Minimal kernel

```
1 typedef unsigned char uint8_t;
2 typedef unsigned int uint32_t;
3 typedef uint32_t size_t;
```

5

```
 4
 5 extern char __bss[], __bss_end[], __stack_top[];
 6
 7 void *memset(void *buf, char c, size_t n) {
 8   uint8_t *p = (uint8_t *)buf;
 9   while (n--)
10     *p++ = c;
11   return buf;
12 }
13
14 void kernel_main(void) {
15   memset(__bss, 0, (size_t) __bss_end - (size_t) __bss);
16
17   for(;;);
18 }
19
20 __attribute__((section(".text.boot")))
21 __attribute__((naked))
22 void boot(void) {
23   __asm__ __volatile__(
24     "mv sp, %[stack_top]\n"
25     "j kernel_main\n"
26     :
27     : [stack_top] "r" (__stack_top)
28   );
29 }
```

Listing 6: kernel code

The kernel starts executing from the `boot` function (which was specified as the entry point in the linker script).

1. `__attribute__((section(".text.boot")))` controls the placement of the function in the linker script. Since OpenSBI jumps to `0x80200000` automatically, the boot code needs be to put there

2. `__attribute__((naked))` instructs the compiler not to generate any unnecessary code before and after the function body. This ensures that the inline assembly code is the exact function body.

In the `boot` function, in line (24), the `stack pointer (sp)` is set to the end address of the stack area thats defined in the linker script. Then, it jumps to the `kernel_main` function.

In line (5), `extern` is used because these variables come from the linker script, and not `kernel.c`. The `char` and `[]` that they are addresses, you can take the pointer value of it, and that its type is a pointer to char, which is a byte.

In the `main` function, the `.bss` section is initalized to zero using `memset`. Then, the function enters an infinite loop. This prevents the execution from leaving the controlled code.

3.5   First kernel debugging

Because the kernel enters an infinite loop, there are no indications that the kernel is running correctly. However, you can look at values of registers to see if anything has gone wrong. To do this, open the QEMU monitor and execute the `info registers` command

```
(qemu) info registers

CPU#0
 V       =   0
 pc       80200048
 mhartid  00000000
 mstatus  80006080
 mstatush 00000000
 hstatus  00000000
 vsstatus 00000000
 mip      00000020
 mie      00000008
 mideleg  00001666
 hideleg  00000000
 medeleg  00f0b509
 hedeleg  00000000
 mtvec    800004e0
 stvec    80200000
 vstvec   00000000
 mepc     80200000
 sepc     00000000
 vsepc    00000000
 mcause   00000003
 scause   00000000
 vscause  00000000
 mtval    80010724
 stval    00000000
 htval    00000000
 mtval2   00000000
 mscratch 80046000
 sscratch 00000000
 satp     00000000
 x0/zero  00000000 x1/ra    8000e63e x2/sp    8022004c x3/gp    00000000
 x4/tp    80046000 x5/t0    00000001 x6/t1    00000002 x7/t2    00001000
 x8/s0    80045f40 x9/s1    00000001 x10/a0   8020004c x11/a1   8020004c
 x12/a2   00000000 x13/a3   00000019 x14/a4   00000000 x15/a5   00000001
 x16/a6   00000001 x17/a7   00000005 x18/s2   80200000 x19/s3   00000000
 x20/s4   87e00000 x21/s5   00000000 x22/s6   80006800 x23/s7   00000001
 x24/s8   00002000 x25/s9   80042308 x26/s10  00000000 x27/s11  00000000
 x28/t3   80020ad1 x29/t4   80045f40 x30/t5   0000008c x31/t6   00000000
 fcsr     00000000
 f0/ft0   ffffffff00000000 f1/ft1   ffffffff00000000 f2/ft2   ffffffff00000000 f3/ft3   ffffffff00000000
 f4/ft4   ffffffff00000000 f5/ft5   ffffffff00000000 f6/ft6   ffffffff00000000 f7/ft7   ffffffff00000000
 f8/fs0   ffffffff00000000 f9/fs1   ffffffff00000000 f10/fa0  ffffffff00000000 f11/fa1  ffffffff00000000
 f12/fa2  ffffffff00000000 f13/fa3  ffffffff00000000 f14/fa4  ffffffff00000000 f15/fa5  ffffffff00000000
 f16/fa6  ffffffff00000000 f17/fa7  ffffffff00000000 f18/fs2  ffffffff00000000 f19/fs3  ffffffff00000000
 f20/fs4  ffffffff00000000 f21/fs5  ffffffff00000000 f22/fs6  ffffffff00000000 f23/fs7  ffffffff00000000
 f24/fs8  ffffffff00000000 f25/fs9  ffffffff00000000 f26/fs10 ffffffff00000000 f27/fs11 ffffffff00000000
 f28/ft8  ffffffff00000000 f29/ft9  ffffffff00000000 f30/ft10 ffffffff00000000 f31/ft11 ffffffff00000000
```

4   Hello World!

Instead of having to look at registers to confirm whether something is working or not, want to make it more obvious by outputing a string from the kernel!

4.1   Hello via SBI

To use functions by the SBI, we use the `ecall` function found in line (17) of Listing 8. But before that inline assembly code, we have to ask the compiler to place a few values in specified registers. That is accomplished in lines (7) - (14) in Listing 8. Specifically, the `register` and `__asm__("register name")` does this.

When `ecall` is evoked, the CPU's execution mode switches from kernel mode (S-Mode) to OpenSBI mode (M-Mode) and OpenSBI's processing handler is invoked. Once its finished, it switches back to kernel mode and execution resumes after the `ecall` instruction.

```
1 #pragma once
2
3 struct sbiret {
4   long error;
5   long value;
6 };
```

Listing 7: (kernel.h) using ecall

```
1  #include "kernel.h"
2  #include "common.h"
3
4  extern char __bss[], __bss_end[], __stack_top[];
5
6  struct sbiret sbi_call(long arg0, long arg1, long arg2, long arg3, long arg4, long
        arg5, long fid, long eid) {
7    register long a0 __asm__("a0") = arg0;
8    register long a1 __asm__("a1") = arg1;
9    register long a2 __asm__("a2") = arg2;
10   register long a3 __asm__("a3") = arg3;
11   register long a4 __asm__("a4") = arg4;
12   register long a5 __asm__("a5") = arg5;
13   register long a6 __asm__("a6") = fid;
14   register long a7 __asm__("a7") = eid;
15
16   __asm__ __volatile__(
17     "ecall"
18     : "=r"(a0), "=r"(a1)
19     : "r"(a0), "r"(a1), "r"(a2), "r"(a3), "r"(a4), "r"(a5), "r"(a6), "r"(a7)
20     : "memory"
21   );
22   return (struct sbiret){.error = a0, .value=a1};
23 }
24
25 void putchar(char ch) {
26   sbi_call(ch, 0, 0, 0, 0, 0, 0, 1);
27 }
28
29 void kernel_main(void) {
30   memset(__bss, 0, (size_t) __bss_end - (size_t) __bss);
31
32   printf("\n\nHello %s\n", "World!");
33   for (int i = 0; s[i] != '\0'; i++) {
34     putchar(s[i]);
35   }
36
37   for(;;) {
38     __asm__ __volatile__("wfi");
39   }
40 }
```

Listing 8: (kernel.c) using ecall

In order to write to the console, need to use the OpenSBI call `Console Putchar`, which puts a character to the console. To make the call, we select function 0 and extension 1 (by putting 0 in register `a6` and 1 in register `a7`) [2]. In addition, SBI functions must return a pair of values in `a0` and `a1` where the value in `a0` is the error code. This is why we have the `struct sbiret` definition in Listing 7. In addition, `Console Putchar` is guarenteed not to change registers `a2` to `a7`.

4.2   printf function

Now, we can implement the `printf` function which, under the hood, simply prints the characters supplied to the function (using `Console Putchar`). In this `printf`, we are only going to support 3 format specifiers:

1. `%d` (decimal)

2. `%x` (hexadecimal)

3. `%s` (string)

```
1 #pragma once
2
3 #define va_list __builtin_va_list
4 #define va_start __builtin_va_start
5 #define va_end __builtin_va_end
6 #define va_arg __builtin_va_arg
7
8 void printf(const char *fmt, ...);
```

Listing 9: (common.h) header file for printf

```
1 #include "common.h"
2
3 void putchar(char ch);
4
5 void printf(const char *fmt, ...) {
6   va_list vargs;
7   va_start(vargs, fmt);
8
9   while (*fmt) {
10     if (*fmt == '%') {
11       fmt++;
12       switch (*fmt) {
13         case '\0': // at the end of the format string
14           putchar('%');
15           goto end;
16         case '%':
17           putchar('%');
18           break;
19         case 's': {
20           const char *s = va_arg(vargs, const char *);
21           while (*s) {
22             putchar(*s);
23             s++;
24           }
25           break;
26         }
27         case 'd': {
28           int value = va_arg(vargs, int);
29           unsigned magnitude = value;
30           if (value < 0) {
31             putchar('-');
32             magnitude = -magnitude;
33           }
34
35           unsigned divisor = 1;
36           while (magnitude / divisor > 9)
37             divisor *= 10;
38
39           while (divisor > 0) {
40             putchar('0' + magnitude / divisor);
41             magnitude %= divisor;
42             divisor /= 10;
43           }
44           break;
45         }
46         case 'x': {
47           unsigned value = va_arg(vargs, unsigned);
48           for (int i = 7; i >= 0; i--) {
```

```
49              unsigned nibble = (value >> (i * 4)) & 0xf;
50              putchar("0123456789abcdef"[nibble]);
51          }
52        }
53      }
54    } else {
55      putchar(*fmt);
56    }
57    fmt++;
58  }
59  end:
60    va_end(vargs);
61 }
```

Listing 10: (common.c) C file for printf

References

[1] Alex Bennee. *Anatomy of a Boot, a QEMU perspective*. URL: https://www.qemu.org/2020/07/03/anatomy-of-a-boot/.

[2] Stephen Marz. *OpenSBI Calls*. URL: https://courses.stephenmarz.com/my-courses/cosc562/risc-v/opensbi-calls/.

[3] Seiya Nuta. *OS in 1000 Lines*. URL: https://operating-system-in-1000-lines.vercel.app/en/.